

COMP9024: Data Structures and Algorithms

Priority Queues and Disjoint Set Union-Find Data
Structures

Contents

- Priority queue ADT
- Heap-based priority queues
- Binomial Heaps
- Disjoint set union-find data structures and algorithms

Priority Queue ADT

- A priority queue stores a collection of items.
- Each **item** is a pair (key, value), where key is the priority of the item.
- Main operations of the Priority Queue ADT:
 - `Insert(k, x)`
Inserts an item with key k and value x.
 - `RemoveMin()` (`RemoveMax()`)
Removes and returns the item with smallest key (largest key). We consider `RemoveMin()` only. Implementation of `RemoveMax()` is similar.
- Additional operations
 - `Min()` (`Max()`)
returns, but does not remove, an entry with smallest key (largest key)
 - `Size()`, `IsEmpty()`
- Applications:
 - Standby flyers
 - Auctions
 - Stock market

Total Order Relations

- Keys in a priority queue can be arbitrary objects on which a total order is defined.
- Two distinct entries in a priority queue can have the same key.
- Mathematical concept of total order relation \leq
 - Reflexive property:
 $x \leq x$
 - Antisymmetric property:
 $x \leq y \wedge y \leq x \Rightarrow x = y$
 - Transitive property:
 $x \leq y \wedge y \leq z \Rightarrow x \leq z$

Priority Queue Sorting

- We can use a priority queue to sort a set of comparable elements:
 1. Insert the elements one by one with a series of **Insert** operations.
 2. Remove the elements in sorted order with a series of **RemoveMin** operations.
- The running time of this sorting algorithm depends on the priority queue implementation.

Algorithm *PQ-Sort(S)*

Input sequence S

Output sequence S sorted in non-decreasing order

```
{ Create an empty priority queue  $P$ ;  
  while ( $\neg IsEmpty(S)$ )  
    {  $e = RemoveFirst(S)$ ;  
       $Insert(P, e)$  ;  
    }  
  while ( $\neg IsEmpty(P)$ )  
    {  $e = RemoveMin(P)$ ;  
       $InsertLast(S, e)$ ;  
    }  
}
```

List-based Priority Queue

- Implementation with an unsorted list:



- Performance:
 - **Insert** takes $O(1)$ time since we can insert the item at the beginning or end of the list.
 - **RemoveMin** and **Min** take $O(n)$ time since we have to traverse the entire list to find the smallest key.

- Implementation with a sorted list:



- Performance:
 - **Insert** takes $O(n)$ time since we have to find the place where to insert the item.
 - **RemoveMin** and **Min** take $O(1)$ time, since the smallest key is at the beginning.

Selection-Sort

- Selection-sort is a variation of PQ-sort where the priority queue is implemented with an unsorted list.
- Running time of Selection-sort:
 1. Inserting the elements into the priority queue with n insert operations takes $O(n)$ time.
 2. Removing the elements in sorted order from the priority queue with n RemoveMin operations takes time proportional to
$$1 + 2 + \dots + n$$
- Selection-sort runs in $O(n^2)$ time.

Selection-Sort Example

	<i>List S</i>	<i>Priority Queue P</i>
Input:	(7,4,8,2,5,3,9)	()
Phase 1		
(a)	(4,8,2,5,3,9)	(7)
(b)	(8,2,5,3,9)	(7,4)
..
.	.	.
(g)	()	(7,4,8,2,5,3,9)
Phase 2		
(a)	(2)	(7,4,8,5,3,9)
(b)	(2,3)	(7,4,8,5,9)
(c)	(2,3,4)	(7,8,5,9)
(d)	(2,3,4,5)	(7,8,9)
(e)	(2,3,4,5,7)	(8,9)
(f)	(2,3,4,5,7,8)	(9)
(g)	(2,3,4,5,7,8,9)	()

Insertion-Sort


- Insertion-sort is the variation of PQ-sort where the priority queue is implemented with a sorted list.
- Running time of Insertion-sort:
 1. Inserting the elements into the priority queue with n **Insert** operations takes time proportional to $1 + 2 + \dots + n$
 2. Removing the elements in sorted order from the priority queue with a series of n **RemoveMin** operations takes $O(n)$ time.
- Insertion-sort runs in $O(n^2)$ time.

Insertion-Sort Example


	<i>List S</i>	<i>Priority queue P</i>
Input:	(7,4,8,2,5,3,9)	()
Phase 1		
(a)	(4,8,2,5,3,9)	(7)
(b)	(8,2,5,3,9)	(4,7)
(c)	(2,5,3,9)	(4,7,8)
(d)	(5,3,9)	(2,4,7,8)
(e)	(3,9)	(2,4,5,7,8)
(f)	(9)	(2,3,4,5,7,8)
(g)	()	(2,3,4,5,7,8,9)
Phase 2		
(a)	(2)	(3,4,5,7,8,9)
(b)	(2,3)	(4,5,7,8,9)
..
.	.	.
(g)	(2,3,4,5,7,8,9)	()

In-place Insertion-sort

- Instead of using an external data structure, we can implement selection-sort and insertion-sort in-place.
- A portion of the input list itself serves as the priority queue.
- For in-place insertion-sort
 - We keep sorted the initial portion of the list.
 - We can use **swaps** instead of modifying the list.




Heaps



Heaps

- A min-heap is a binary tree storing keys at its nodes and satisfying the following properties:
 - **Heap-Order:** for every node v other than the root, $key(v) \geq key(parent(v))$
 - **Complete Binary Tree:** let h be the height of the heap
 - for $i = 0, \dots, h - 1$, there are 2^i nodes of depth i
 - at depth h , all the nodes are as far left as possible
- The last node of a heap is the rightmost node of depth h .

- A max-heap satisfies a different heap order property:
 - For every node v other than the root, $key(v) \leq key(parent(v))$
- We consider min-heap only




Height of a Heap

- **Theorem:** A heap storing n keys has height $O(\log n)$.


Proof: (we apply the complete binary tree property)

- Let h be the height of a heap storing n keys.
- Since there are 2^i keys at depth $i = 0, \dots, h-1$ and at least one key at depth h , we have $n \geq 1 + 2 + 4 + \dots + 2^{h-1} + 1$.
- Thus, $n \geq 2^h$, i.e., $h \leq \log n$.





Heaps and Priority Queues

- We can use a heap to implement a priority queue.
- We store a (key, element) item at each node.
- We keep track of the position of the last node.
- For simplicity, we show only the keys in the pictures.





Insertion into a Heap

- Operation **Insert** of the priority queue ADT corresponds to the insertion of a key k to the heap.
- The insertion algorithm consists of three steps:
 - Find the insertion node z (the new last node)
 - Store k at z
 - Restore the heap-order property (discussed next)





Upheap

- After the insertion of a new key k , the heap-order property may be violated.
- Algorithm upheap restores the heap-order property by swapping k along an upward path from the insertion node.
- Upheap terminates when the key k reaches the root or a node whose parent has a key smaller than or equal to k .
- Since a heap has height $O(\log n)$, upheap runs in $O(\log n)$ time.





Removal from a Heap

- Method RemoveMin of the priority queue ADT corresponds to the removal of the root key from the heap.
- The removal algorithm consists of three steps
 - Replace the root key with the key of the last node w
 - Remove w
 - Restore the heap-order property (discussed next)




Downheap

- After replacing the root key with the key k of the last node, the heap-order property may be violated.
- Algorithm downheap restores the heap-order property by swapping key k along a downward path from the root.
- Downheap terminates when key k reaches a leaf or a node whose children have keys greater than or equal to k .
- Since a heap has height $O(\log n)$, downheap runs in $O(\log n)$ time.



Updating the Last Node

- The insertion node can be found by traversing a path of $O(\log n)$ nodes:
 - Go up until a left child or the root is reached
 - If a left child is reached, go to the right child
 - Go down left until the next node is null.
- Similar algorithm for updating the last node after a removal.





Heap-Sort

- Consider a priority queue with n items implemented by means of a heap
 - The space used is $O(n)$
 - Operations **Insert** and **RemoveMin** take $O(\log n)$ time
 - Operations **Size**, **IsEmpty**, and **Min** take $O(1)$ time
- Using a heap-based priority queue, we can sort a sequence of n items in $O(n \log n)$ time.
- The resulting algorithm is called heap-sort.
- Heap-sort is much faster than quadratic sorting algorithms, such as insertion-sort and selection-sort.



Array-based Heap Implementation

- We can represent a heap with n keys by means of an array of length $n + 1$.
- For the node at rank i
 - the left child is at rank $2i$
 - the right child is at rank $2i + 1$
- Links between nodes are not explicitly stored.
- The cell of at rank 0 is not used.
- Operation **Insert** corresponds to inserting at rank $n + 1$.
- Operation **RemoveMin** corresponds to removing at rank n .
- Yields in-place heap-sort.




Merging Two Heaps

- We are given two two heaps and a key k .
- We create a new heap with the root node storing k and with the two heaps as subtrees
- We perform downheap to restore the heap-order property.





Bottom-up Heap Construction


- We can construct a heap storing n given keys in using a bottom-up construction with $\log n$ phases.
- In phase i , pairs of heaps with $2^i - 1$ keys are merged into heaps with $2^{i+1} - 1$ keys.




Example (1/4)




Example (2/4)



Example (3/4)




Example (4/4)





Analysis

- We visualize the worst-case time of a downheap with a proxy path that goes first right and then repeatedly goes left until the bottom of the heap (this path may differ from the actual downheap path)
- Since each node is traversed by at most two proxy paths, the total number of nodes of the proxy paths is $O(n)$.
- Thus, bottom-up heap construction runs in $O(n)$ time.
- Bottom-up heap construction is faster than n successive insertions and speeds up the first phase of heap-sort.



Binomial Heaps



Merge Two Heaps with Different Sizes

Merge(H1,H2): Merge two heaps H1 and H2 with sizes m and n.

Algorithm 1: Insert each entry from H1 and H2 into a new heap:

Running time: $O((m+n) \log (m+n))$



Algorithm 2: Use the bottom-up heap construction algorithm

Running Time: $O(m+n)$

Can we merge two heaps in $O(\log(m+n))$ time?

Binomial Trees

- Recursive Definition of Binomial tree B_k of height k :
 - B_0 = single root node
 - B_k = Attach B_{k-1} to root of another B_{k-1}




Properties of Binomial Trees

For the binomial tree B_k :

1. There are 2^k nodes
2. The height of B_k is k
3. There are exactly $\binom{k}{i}$ (binomial coefficient) nodes at depth i for $i = 0, 1, \dots, k$


Binomial Heaps

- A binomial heap H_k is a set of binomial trees B_0, B_1, \dots, B_k where each binomial tree is heap-ordered:
 - The key of each node \geq the key of the parent
- The root of each binomial tree in H_k contains the smallest key in that tree.



findMin()

- Traverse all the roots, taking $O(\log n)$ time




Merge Two Binomial Heaps (1/6)


- Key ideas: merge individual pairs of heaps with the same height
- Steps for merging two binomial heaps:
 1. Create a new empty binomial heap
 2. Start with B_k for the smallest k
 3. If there is only one B_k , add B_k to the new binomial heap and go to Step 3 with $k = k + 1$
 4. Merge two B_k 's into a new B_{k+1} by making the root with a larger key the child of the other root. Go to Step 3 with $k = k + 1$
- Time complexity: $O(\log (m+n))$, where m and n are the sizes of two heaps.

Merge Two Binomial Heaps (2/6)

H1:




H2:




Merge Two Binomial Heaps (3/6)

H1:




H2:




Merge Two Binomial Heaps (4/6)

H1:




H2:




Merge Two Binomial Heaps (5/6)


H1:



H2:



Merge Two Binomial Heaps (6/6)



Insertion

- Create a single node tree B_0 with the new item and merge with the existing heap
- Time complexity: $O(\log n)$

How Many Binomial Trees in a Binomial Heap?

Consider a binomial heap with n nodes

- Convert n into a binary number $b_k b_{k-1} \dots b_1 b_0$
- If $b_i \neq 0$ ($i=0, 1, \dots, k$), the binomial tree B_i is not empty

Example: $n=28$.

- $28=16+8+4=11100$. So $k=4$. The binomial heap consists of B_4 (16 nodes), B_3 (8 nodes) and B_2 (4 nodes).

removeMin() (1/4)


Steps:

1. Find the tree B_k with the **smallest** root
2. Remove B_k from the heap
3. Keep the entry stored at the root of B_k (return value) and remove the root of B_k (now we have a new forest B_0, B_1, \dots, B_{k-1})
4. Merge this new forest with remainder of the original
5. Return the entry with the min key


Run time analysis:

- Step 1 is $O(\log n)$, Step 2 and Step 3 are $O(1)$, and Step 4 is $O(\log n)$
- Total time complexity is $O(\log n)$


removeMin() (2/4)



removeMin() (3/4)



removeMin() (4/4)



Disjoint Set Union-Find Structures

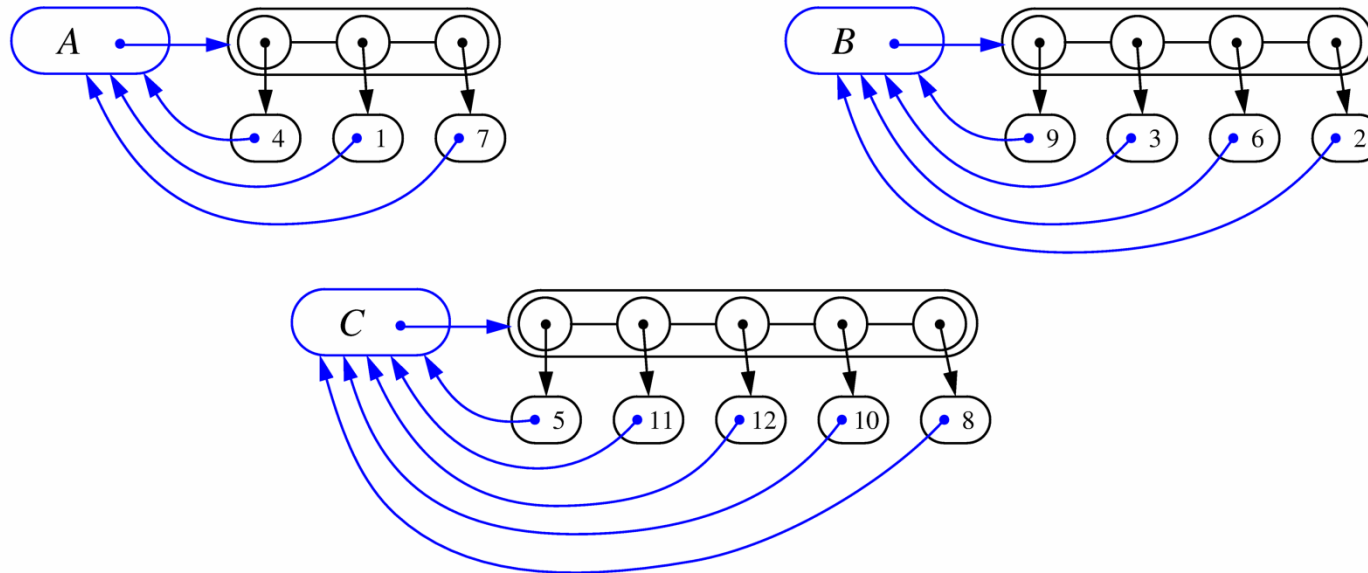


Disjoint Set Union-Find Operations

- **MakeSet(x)**: Create a singleton set containing the element x and return the position storing x in this set.
- **Union(A, B)**: Return the set $A \cup B$, destroying the old A and B .
- **Find(e)**: Return the set containing the element e .

List-based Implementation

- Each set is stored in a sequence represented with a linked-list
- Each node should store an object containing the element and a reference to the set name




Analysis of List-based Representation

- When doing a union, always move elements from the smaller set to the larger set
 - Each time an element is moved it goes to a set of size at least double its old set
 - Thus, an element can be moved at most $O(\log n)$ times
- Total time needed to do n unions and finds is $O(n \log n)$.



Tree-based Implementation

- Each element is stored in a node, which contains a pointer to a **set** name
- A node v whose set pointer points back to v is also a set name
- Each set is a tree, rooted at a node with a self-referencing set pointer
- For example: The sets “1”, “2”, and “5”:




Union-Find Operations

- To do a **Union**, simply make the root of one tree point to the root of the other
- To do a **Find**, follow set-name pointers from the starting node until reaching a node whose set-name pointer refers back to itself



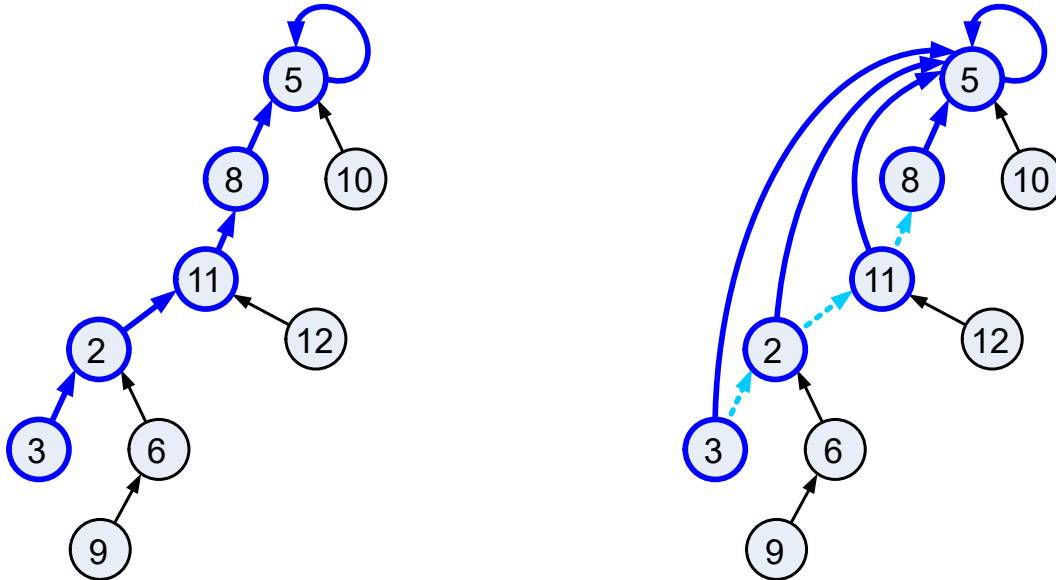
Union-Find Heuristic 1

- Union by size:
 - When performing a union, make the root of smaller tree point to the root of the larger
- Implies $O(n \log n)$ time for performing n union-find operations:
 - Each time we follow a pointer, we are going to a subtree of size at least double the size of the previous subtree
 - Thus, we will follow at most $O(\log n)$ pointers for any find.



Union-Find Heuristic 2

- Path compression:
 - After performing a find, compress all the pointers on the path just traversed so that they all point to the root



- Implies $O(n \log^* n)$ time for performing n union-find operations:
 - Proof is somewhat involved.

Proof of $\log^* n$ Amortized Time

- For each node v that is a root
 - define $n(v)$ to be the size of the subtree rooted at v (including v)
 - identified a set with the root of its associated tree.
- We update the size field of v each time a set is unioned into v . Thus, if v is not a root, then $n(v)$ is the largest the subtree rooted at v can be, which occurs just before we union v into some other node whose size is at least as large as v 's.
- For any node v , then, define the **rank** of v , which we denote as $r(v)$, as $r(v) = \lfloor \log n(v) \rfloor$:
- Thus, $n(v) \geq 2^{r(v)}$.
- Also, since there are at most n nodes in the tree of v , $r(v) = \lfloor \log n \rfloor$, for each node v .

Proof of $\log^* n$ Amortized Time (2)

- For each node v with parent w :
 - $r(w) > r(v)$
- **Claim:** *There are at most $n/2^s$ nodes of rank s .*
- **Proof:**
 - Since $r(v) < r(w)$, for any node v with parent w , ranks are monotonically increasing as we follow parent pointers up any tree.
 - Thus, if $r(v) = r(w)$ for two nodes v and w , then the nodes counted in $n(v)$ must be separate and distinct from the nodes counted in $n(w)$.
 - If a node v is of rank s , then $n(v) \geq 2^s$.
 - Therefore, since there are at most n nodes total, there can be at most $n/2^s$ that are of rank s .

Proof of $\log^* n$ Amortized Time (3)

- Definition: Tower of two's function:
 - $t(i) = 2^{t(i-1)}$
- Nodes v and u are in the same rank group g if
 - $g = \log^*(r(v)) = \log^*(r(u))$:
- Since the largest rank is $\log n$, the largest rank group is
 - $\log^*(\log n) = (\log^* n) - 1$

Proof of $\log^* n$ Amortized Time (4)

- Charge 1 cyber-dollar per pointer hop during a find:
 - If w is the root or if w is in a different rank group than v , then charge the find operation one cyber-dollar.
 - Otherwise (w is not a root and v and w are in the same rank group), charge the node v one cyber-dollar.
- Since there are most $(\log^* n)-1$ rank groups, this rule guarantees that any find operation is charged at most $\log^* n$ cyber-dollars.

Proof of $\log^* n$ Amortized Time (5)

- After we charge a node v then v will get a new parent, which is a node higher up in v 's tree.
- The rank of v 's new parent will be greater than the rank of v 's old parent w .
- Thus, any node v can be charged at most the number of different ranks that are in v 's rank group.
- If v is in rank group $g > 0$, then v can be charged at most $t(g) - t(g-1)$ times before v has a parent in a higher rank group (and from that point on, v will never be charged again). In other words, the total number, C , of cyber-dollars that can ever be charged to nodes can be bound as

$$C \leq \sum_{g=1}^{\log^* n - 1} n(g) \cdot (t(g) - t(g-1))$$

Proof of $\log^* n$ Amortized Time (end)

- Bounding $n(g)$:

$$\begin{aligned} n(g) &\leq \sum_{s=t(g-1)+1}^{t(g)} \frac{n}{2^s} \\ &= \frac{n}{2^{t(g-1)+1}} \sum_{s=0}^{t(g)-t(g-1)-1} \frac{1}{2^s} \\ &< \frac{n}{2^{t(g-1)+1}} \cdot 2 \\ &= \frac{n}{2^{t(g-1)}} \\ &= \frac{n}{t(g)} \end{aligned}$$

- Returning to C:

$$\begin{aligned} C &< \sum_{g=1}^{\log^* n - 1} \frac{n}{t(g)} \cdot (t(g) - t(g-1)) \\ &\leq \sum_{g=1}^{\log^* n - 1} \frac{n}{t(g)} \cdot t(g) \\ &= \sum_{g=1}^{\log^* n - 1} n \\ &\leq n \log^* n \end{aligned}$$

Summary

- Priority queue ADT
- List-based priority queues
- Heap-based priority queues
- Bottom-up heap construction
- Binomial heaps
- Disjoint set union-find data structures and algorithms
- Suggested reading:
 - Sedgewick, Ch. 1.3, 9.